

Silicon valley real-estate project

Maxime LU

A dark blue diagonal gradient bar that starts from the bottom left corner and extends towards the top right corner, covering the lower half of the slide.

Sommaire

- Premier jour: Exploration
- Crash test
- Dernier jour: Phase finale

Exploration: Premier jour

Préparation et nettoyage

12/01 début du projet

préparation sur notion

création du repo git

commencement de l'analyse et nettoyage du dataset

détection de Nan et commencement du kaggle sur les missing values

remplacement des Nan par des 0

plotting pour l'EDA

finition de l'EDA

Planning projet immo

<https://simplonline.co/briefs/8830c87-e014-43e3-bc55-d87744ee33e2>

Vue « tableau kanban » ▾

Filtrer Trier Q % ...

Nouvelle page ▾

Projet immobilier ...

À faire 3

En cours 2

Terminé 1

Faire les kaggles

notebook EDA

création repo git

git - faire le README

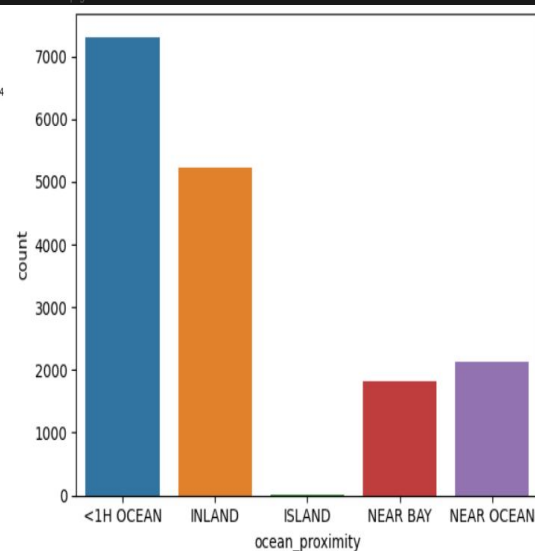
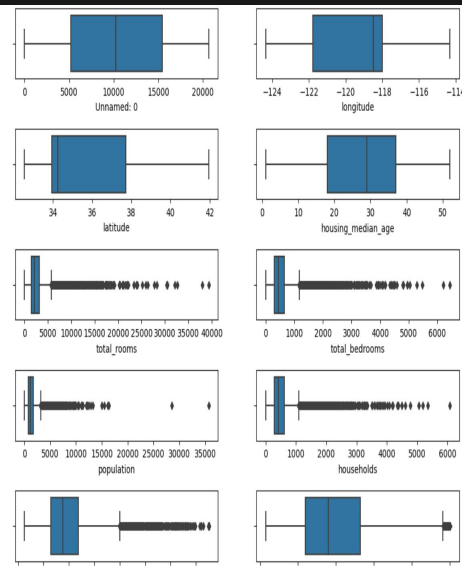
nettoyer le dataset

+ Nouvelle page

ENVOYER LE LIEN GIT SUR
SIMPLONLINE POUR ARTURO

+ Nouvelle page



+ Nouvelle page



Exploration:

Deuxième jour

Notebook inférence



Exploration:

Deuxième jour

Premiers models

13/01

push sur git

création data directory (j'ai mis un placeholder dedans pour le push vide)

création et personalisation de .gitignore

exportation de l'EDA

commencement de l'inférence

test de shapiro et levene pour l'anova

test de pearson

création du notebook model

création du premier model BM (score de 0,006)

test du score

ont peut voir le problème, je n'ai pas scale les donnée, c'est pour que sa mettait 10 a run.

je vais maintenant commencer le kaggle et recommencer le cycle a l'EDA

finition du tp model evaluation

2eme itération de l'EDA

plot de la normalisation et analyse

```
dummy_clf.score(x, y)
```

0.04620881782945736

```
In [5]: log_model = LogisticRegression()
log_model.fit(X_train, y_train)
y_pred = log_model.predict(X_test)
y_prob = log_model.predict_proba(X_test)
log_model.score(X_test, y_test)
```

```
Out[5]: 0.000000000000000000
```

```
In [10]: print(classification_report(y_test, y_pred))
```

| | precision | recall | f1 score | support |
|----------|-----------|--------|----------|---------|
| 22000_0 | 0.00 | 0.00 | 0.00 | 2 |
| 22000_1 | 0.00 | 0.00 | 0.00 | 2 |
| 22000_2 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_3 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_4 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_5 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_6 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_7 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_8 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_9 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_10 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_11 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_12 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_13 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_14 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_15 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_16 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_17 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_18 | 0.00 | 0.00 | 0.00 | 1 |
| 22000_19 | 0.00 | 0.00 | 0.00 | 1 |

second iteration model - score: 0.6350

model iterations (dummy,LinearRegression)

baseline first iteration - score:0.046

Nettoyage: des première iterations

Second iteration - fillna with 0

third iteration - fillna, no outliers

fourth iteration - fill median

fifth iteration - fill median, no outlier

sixth iteration - fill mean

seventh iteration - fill mean, no outlier

Crash test:

création d'itération

eight iteration - encoding

ninth iteration - discretizing, house_median_age

tenth iteration - encoding discretized house_median_age

eleventh iteration - scaling using standard scaler

twelfth iteration - scaling using robust scaler

thirteenth iteration - normalizing

fourteenth iteration - logarithmic scaling (UNFINISHED)

Models:

Vérification des itérations

```
#making a function to do everything i need in a Linear regression
def lineareg(data,target):
    #linear regression
    y = data[target]
    X = data.drop(target, axis=1)
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.3, random_state=1)
    model = LinearRegression()
    model.fit(X_train, y_train)
    print(f'initial model score is {model.score(X_test, y_test)}')
    #cross validation
    K = []
    total_time = []
    score = []
    for k in range(2,20):
        cross_val_results = cross_validate(model, X, y, cv=k)
        total_time.append(sum(cross_val_results['fit_time'])+sum(cross_val_results['score_time']))
        K.append(k)
        score.append(cross_val_results['test_score'].mean())
    #wrote these 2 lines to select the best cross validate value
    best_cv = pd.DataFrame({'K': K, 'score': score})
    cross = best_cv.query(f'score==(max(score))')[['K']].item()
    cv_results = cross_validate(model,X, y,cv=cross, scoring=('r2', 'neg_root_mean_squared_error','neg_mean_absolute_error'))
    r2 = cv_results['test_r2'].mean()
    rmse = cv_results['test_neg_root_mean_squared_error'].mean()
    print(f'r2: {r2}')
    print(f'rmse: {rmse}')
```

tenth iteration model - score: 0.6498 best*

KNN - wanted to try other models so i see what i can get

KNN function

DecisionTree - same as KNN

random forest - doesn't work at all

Makeshift pipeline: notebook predict

pipeline

```
def cleaning(data):  
    #imputation  
    col_nan = data.columns[data.isna().any(0)]  
    for i in col_nan:  
        data = data.fillna(data[i].mean())  
    #discretizing  
    data['binned_age']=pd.cut(data['housing_median_age'],  
                              bins = [0,10,20,30,40,50,np.inf],  
                              labels = ['[0-10]', '[10-20]', '[20-30]', '[30-40]', '[40-50]', '[50+]'])  
    data = data.drop(['housing_median_age'],axis=1)  
    #encoding  
    col_qual = data.select_dtypes(exclude=[np.number])  
    col_qual = col_qual.columns  
    for i in col_qual:  
        encoder = OneHotEncoder()  
        onehot = encoder.fit_transform(data[[i]])  
        onehot = onehot.toarray()  
        onehot_df = pd.DataFrame(onehot, columns=encoder.get_feature_names_out([i]))  
        data = pd.concat([data,onehot_df], axis=1)  
    for i in col_qual:  
        data = data.drop([i],axis=1)  
    return data
```

```
def pickle_import(func):  
    with open(f'{func}.pkl', 'rb') as file:  
        fonction = pickle.load(file)  
    return fonction
```

Git:

finalisation

README.md

projet_immo_silicon

ma contribution au projet: <https://simplonline.co/briefs/8830c8f7-e014-43e3-bc55-d87744ee33e2>

lien de la présentation google slide:

<https://docs.google.com/presentation/d/1DmBTrT6Ni1zlvnbv705Zeyltps2tpSLCOV0gNRxVd6l/edit?usp=sharing>



mungleh Merge pull request #15 from mungleh/tidying

00c58af 2 hours ago 35 commits

| | | |
|----------------------------|------------------|-------------|
| 📁 .ipynb_checkpoints | some changes | 2 hours ago |
| 📁 data | pickled data | last week |
| 📄 .gitignore | ignore | 3 weeks ago |
| 📄 README.md | Update README.md | 3 hours ago |
| 📄 cleaning.pkl | predict notebook | 3 hours ago |
| 📄 model.pkl | predict notebook | 3 hours ago |
| 📄 notebook_EDA.ipynb | predict notebook | 3 hours ago |
| 📄 notebook_inference.ipynb | some changes | 2 hours ago |
| 📄 notebook_model.ipynb | predict notebook | 3 hours ago |
| 📄 notebook_predict.ipynb | predict notebook | 3 hours ago |

mungleh

